

Introduction to Database Systems: CS312

Constraints and Integrity Constraints

Oliver Bonham-Carter

4 Feb 2019

Integrity Constraints

Integrity Constraints

Nulls

Data Types

Affinity Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James Bond

Consider this

- The CONSTRAINTS are an integrity which defines some conditions that restrict the column to contain the true data while inserting or updating or deleting.
- Integrity constraints provide a mechanism for ensuring that data conforms to guidelines specified by the database administrator. The most common types of constraints include:
 - UNIQUE constraints: To ensure that a given column is unique
 - NOT NULL constraints: To ensure that no null values are allowed
 - FOREIGN KEY constraints: To ensure that two keys share a primary key to foreign key relationship
 - Ensure that a link exists between two tables.

Integrity Constraints

Nulls

Data Types

Affinity Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James Bond

Consider this

- Constraints can be used for these purposes in a data warehouse:
 - Data cleanliness
 - Constraints verify that the data in the data warehouse conforms to a basic level of data consistency and correctness, preventing the introduction of dirty data.
 - Query optimization
 - Although constraints can be useful in many aspects of query optimization, constraints are particularly important for query rewrite of materialized views.

Simple NULL constraint demo

Integrity
Constraints

Nulls

Data Types
Affinity
Constraints
Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

Spot the integrity constraint's influence

```
/*Create table*/  
drop table company;  
create table company(  
    Id text NOT NULL,  
    Name text NOT NULL);
```

```
/*Good insert command: complete tuple allowed*/  
INSERT INTO company VALUES("COM1","T S LTD.");  
SELECT * FROM company;
```

```
/*Good insert command: Empty spaces are allowed*/  
INSERT INTO company VALUES("COM1","");  
SELECT * FROM company;
```

```
/*Bad insert command: NULL is not allowed*/  
INSERT INTO company VALUES("COM1",NULL);
```

Constraints

Define variables by data type!

Integrity
Constraints

Nulls

Data Types

Affinity
Constraints
Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

- **int**: Integer (a finite subset of the integers that is machine dependent).
- **smallint**: Small integer (a machine dependent subset of the integer domain type).
- **numeric(p,n)**: Fixed point number, with user-specified precision of p digits, with n digits to the right of decimal point. (This allows for number comparisons using operators.)
- **real, double precision**: Floating point and double precision floating point numbers, with machine dependent precision.
- **float(n)**: Floating point number, with user specified precision of at least n digits.
- **NOT NULL**: Ensure that a value is placed here or reject the insertion.

Constraints

Define variables by data type!

Integrity
Constraints

Nulls

Data Types

Affinity
Constraints
Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

- **char(*n*)**: Fixed length character string, with user-specified length *n*.
 - Used to store character string value of fixed length
 - The maximum num of chars (not important to SQLite3)
 - About 50 per cent faster than VARCHAR
- **varchar(*n*)**: Variable length character strings, with user specified maximum length *n*.
 - Used to store variable length alphanumeric data
 - The maximum num of chars (not important to SQLite3)
 - Slower than CHAR

Affinity Constraints

A small subset of accepted data types that SQLite will accept

Integrity
Constraints

Nulls

Data Types

Affinity
Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

| Example Typenames From The CREATE TABLE Statement or CAST Expression | Resulting Affinity | Rule Used To Determine Affinity |
|---|--------------------|---------------------------------|
| INT INTEGER TINYINT SMALLINT MEDIUMINT BIGINT UNSIGNED BIG INT INT2 INT8 | INTEGER | 1 |
| CHARACTER(20) VARCHAR(255) VARYING CHARACTER(255) NCHAR(55) NATIVE CHARACTER(70) NVARCHAR(100) TEXT CLOB | TEXT | 2 |
| BLOB <i>no datatype specified</i> | BLOB | 3 |
| REAL DOUBLE DOUBLE PRECISION FLOAT | REAL | 4 |
| NUMERIC DECIMAL(10,5) BOOLEAN DATE DATETIME | NUMERIC | 5 |

Adding Constraints to Create

Integrity
Constraints

Nulls

Data Types

Affinity
Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

```
CREATE TABLE relationshipTable (  
    A1 D1,  
    A2 D2,  
    ...,  
    An Dn,  
    (integrity-constraint1),  
    ...,  
    (integrity-constraintk));
```

- `relationshipTable` is the name of the relationship
- Each A_i is an attribute name in the schema of relation `relationshipTable`
- D_i is the data type of values in the domain of attribute A_i
 - The D_i constrains the particular type of entry

Unique Constraint in the Code

Integrity
Constraints

Nulls
Data Types
Affinity
Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

```
/*Two constraints?*/  
DROP TABLE instructor;  
CREATE TABLE instructor (  
    ID CHAR UNIQUE,  
    name VARCHAR NOT NULL,  
    dept_name VARCHAR,  
    salary VARCHAR  
);
```

```
/******PSSST! Now Add some department information *****/  
insert into instructor values ('10211', 'Smith', 'Biology', 66000);  
  
/*Any trouble inserting this next line?*/  
insert into instructor values ('10212', null, 'Biology', 66000);  
insert into instructor values ('10211', 'Franklin', 'Biology', 66000);
```

- NULL and repeating UNIQUE values are not inserted



Defining a Table with Primary a Key?

Integrity
Constraints

Nulls

Data Types

Affinity
Constraints

Pseudocode

Primary Keys

AgentsDB

Bond, James
Bond

Consider this

ID is unique, Salary bound by numbers

```
/*Two constraints?*/  
DROP TABLE Employee;  
CREATE TABLE Employee (  
    ID CHAR PRIMARY KEY,  
    name VARCHAR NOT NULL,  
    dept_name VARCHAR,  
    salary NUMERIC(8,2)  
);
```

```
/******PSSST! Now Add some secret information *****/  
INSERT INTO Employee VALUES("001","Jimmy", "secretService", 1000000);  
INSERT INTO Employee VALUES("002","Stevie", "secretService", 1000000);  
INSERT INTO Employee VALUES("003","Frankie", "secretService", 10);  
INSERT INTO Employee VALUES("004","Robbie", "secretService", 10A);  
/* Oops! Robbie's salary information has a typographical error*/  
INSERT INTO Employee VALUES("004","Robbie", "secretService", 100);  
INSERT INTO Employee VALUES("004","Jamie", "secretService", 500);  
/* Error: UNIQUE constraint failed: Employee.ID */  
/* Huh?! */
```

Integrity Constraints in CREATE TABLE

Integrity
Constraints

Primary Keys

AgentsDB

Bond, James
Bond

Consider this



- NOT NULL (*but you already knew that!*)
- **Primary Keys:** A primary key is a column or group of columns used to identify the uniqueness of rows in a table.
 - Each table has one and only one primary key.
- **Foreign Keys:** A column (or columns) that references a column (most often the primary key) of another table.
 - The purpose of the foreign key is to ensure referential integrity of the data. In other words, only values that are supposed to appear in the database are permitted.

AgentsDB: Two Tables, One With a Primary Key

Integrity
Constraints

Primary Keys

AgentsDB

Bond, James
Bond

Consider this



File: agentsDB_build.txt

Two Tables, One with Primary Key

```
/* Accepts no redundancy */  
DROP TABLE Agents1;  
CREATE TABLE Agents1  
( last_name VARCHAR NOT NULL,  
  first_name VARCHAR NOT NULL,  
  address VARCHAR,  
  CONSTRAINT agents_pk  
    PRIMARY KEY (last_name, first_name)  
);  
  
/* Accepts redundancy */  
DROP TABLE Agents2;  
CREATE TABLE Agents2  
( last_name VARCHAR NOT NULL,  
  first_name VARCHAR NOT NULL,  
  address VARCHAR  
);
```

Try Your Insert Twice

Let's try

Integrity
Constraints

Primary Keys

AgentsDB

Bond, James
Bond

Consider this



- Insert agent names into both tables.
- Try the same INSERT commands again!
- Which commands work?

```
INSERT INTO Agents1 VALUES ("Bond","James","123 abc street");  
INSERT INTO Agents2 VALUES ("Bond","James","123 abc street");
```

How is the database set up?

- .tables (The tables are of the DB)
- .schema (How the data is stored in the tables)

What data is stored in each table?

- select * from agentsConst;
- select * from agents;
- (note the '*' for the column wildcard)

What Ever happened to “James Bond”?

Integrity
Constraints

Primary Keys

AgentsDB

Bond, James
Bond

Consider this



- There can only be **one** “James Bond”
- The name “James Bond” could not be inserted more than once in our base
- Constraints were in place to ensure *distinguishable* rows



THINK

- Can you build a new database table with two (or more) types of constraints?
- For instance, try to alter an earlier database for which you have the *build* file to recreate it (in case anything goes *dreadfully* wrong)