



Introduction to Database Systems: CS312

Cassandra Theory

Oliver Bonham-Carter

8 April 2019

A NoSQL Database Management System



- Apache Cassandra is a massively scalable open source non-relational database
- Offers continuous availability, linear scale performance, operational simplicity and easy data distribution across multiple data centers and cloud availability zones.
- Founded at Facebook 2008, developed at Apache in 2010

- <http://cassandra.apache.org/>
- <https://academy.datastax.com/planet-cassandra/cassandra>

Cassandra is different from relational DBs

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Relational (Sqlite3)	Cassandra
Handles moderate incoming data velocity	Handles high incoming data velocity
Data arriving from one/few locations	Data arriving from many locations
Manages primarily structured data	Manages all types of data
Supports complex/nested transactions	Supports simple transactions
Single points of failure with failover	No single points of failure; constant uptime
Supports moderate data volumes	Supports very high data volumes

Cassandra is different from relational DBs

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Relational (Sqlite3)	Cassandra
Centralized deployments	Decentralized deployments
Data written in mostly one location	Data written in many locations
Supports read scalability (with consistency sacrifices)	Supports read and write scalability
Deployed in vertical scale up fashion	Deployed in horizontal scale out fashion

Some of its users ...

... and others, too!

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces


Inserting
Data

Shutting
Down

Consider this




Key Features



- Built-for-scale architecture: Cassandra is capable of handling petabytes of information and thousands of concurrent users/operations per second (across multiple data centers) as easily as it can manage much smaller amounts of data and user traffic.
- Unlike other master-slave or sharded systems, Cassandra has no single point of failure and therefore is capable of offering true continuous availability.

Key Features


Distributed and Decentralized



- Distributed: able to be run across several machines at diverse locations
- Active everywhere design: all nodes may be written to and read from.
- No master-slave configurations: all nodes *gossip* meaning they share information using peer-to-peer architecture

Key Features


Elastic Scalability



- Horizontal scaling: adding more machines to handle loads (a linear increase in performance)
- Flexible and dynamic data model: supports modern data types with fast writes and reads.

Key Features


High Availability and Fault Tolerance



- No single point of failure: no “master” node
- Continuous availability: offers redundancy of both data and node function, which eliminate single points of failure and provide constant uptime.

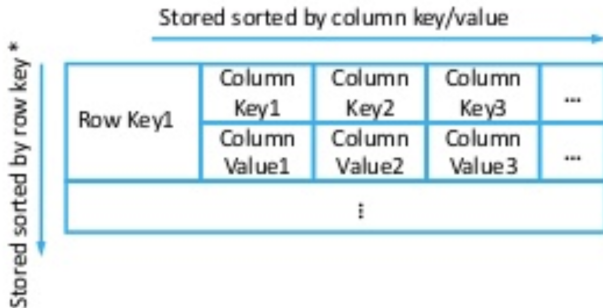
Key Features

Tune up / down the consistency factor



- Tunable Consistency: chose between *strong* and *eventual* consistency
- Adjustable read and write operations (separately)
- Conflicts are solved during reading while the focus is on the write performance.

Tables and Columns



- Data is stored in sparse multidimensional hash tables
- Rows may have multiple columns, not necessary to have same number of columns for each row
- No relations as in Sqlite3
- Each row has unique key that also serves partitioning

Database Language Guide

SQL systems versus CQL (i.e., Cassandra's Query Client)

SQL	CQL	Elaboration
Database	Keyspace	Contain tables. A <i>keyspace</i> defines the replication factor (i.e., the number of replica nodes for ensuring reliability and fault tolerance) and replication strategy for all tables that it contains.
Materialized view	Table + Partition	A CQL table defines a schema much like an SQL table. However, CQL tables contain partitions and each partition contains rows. The combination of a CQL table plus a partition is similar to a materialized view in SQL. A macro-like system of running queries.

Database Language Guide

SQL systems versus CQL (i.e., Cassandra's Query Client)

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

SQL	CQL	Elaboration
Primary key		An SQL primary key is a unique identifier per row. There is no direct equivalent in CQL, although the term "primary key" is used in CQL.
	Primary key	A CQL primary key is a composite key that may define the partition key and optionally clustering columns.
Column	Column	The concept of a column is very similar in Cassandra vs. an RDBMS. Although how a column is physically stored is very different in Cassandra vs. an RDBMS.
Value	Value	The concept of a value is very similar in Cassandra vs. an RDBMS.

Database Language Guide

SQL systems versus CQL (i.e., Cassandra's Query Client)

SQL	CQL	Elaboration
ORDER BY	Clustering columns	Cassandra stores data in sorted order. Therefore, you achieve the equivalent of an SQL ORDER BY through the selection of clustering columns.
JOIN	Achieved via materialized view	As mentioned above, a CQL table plus partition is conceptually closer to a materialized view than a relational table. In a materialized view in an RDBMS you would achieve the equivalent of a JOIN by denormalizing data. The same concept applies to Cassandra where you denormalize data.

Ref: <http://exponential.io/blog/2015/01/08/cassandra-terminology/>

Getting started with Cassandra

- Copy your Cassandra setup tar file (apache-cassandra-3.11.2-bin.tar.gz) to a desktop directory. Do not uncompress this file in your submission directory!!!
- Click on this file to unpack its contents
- Locate the bin directory and then locate file: `cassandra` and file: `cqlsh`

Lab PCs: Setup Java 8 in the path from the terminal

```
JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/
```

Start the cassandra server
Note: Control-C to exit.

```
./cassandra -f
```

With new another terminal, start the `cqlsh` client

```
./cqlsh
```

Keyspaces

Similar to a schema

Find keyspaces (i.e., resembling the schema concept of SQLite3 systems or tables)

```
describe keyspaces;  
describe tables;
```

Start a new keyspaces

```
create keyspace mydb with replication =  
{ 'class':'SimpleStrategy',  
  'replication_factor':1 };
```

Use a keyspace

```
use mydb;
```

Remove a keyspace

```
/*Drop the "mydb" keyspace*/  
DROP KEYSPACE mydb;
```


Inserting

Very Similar to SQL

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Creating
Indexes

Shutting
Down

Consider this

Build a Table

```
create table emp(  
    empid int primary key,  
    emp_first varchar,  
    emp_last varchar,  
    emp_dept varchar);
```

Insert Data Using CQL INSERT command

```
insert into emp  
    ( empid, emp_first, emp_last, emp_dept )  
    values ( 1, 'Fred', 'Smith', 'English' );  
insert into emp  
    ( empid, emp_first, emp_last, emp_dept )  
    values ( 2, 'Bob', 'Alison', 'English' );  
insert into emp  
    ( empid, emp_first, emp_last, emp_dept )  
    values ( 3, 'Judy', 'Miller', 'French' );  
insert into emp  
    ( empid, emp_first, emp_last, emp_dept )  
    values ( 4, 'Jasmin', 'Jones', 'Computer Science' );
```

General Query

Study the schema and indexes (*query-able* columns)

```
describe schema
```

Simple Query

```
select * from emp;
```

```
cqlsh:mydb> select * from emp;
```

empid	emp_dept	emp_first	emp_last
5	French	Megan	Douglas
1	English	Fred	Smith
8	Computer Science	Monroe	Monderson
2	English	Bob	Alison
4	Computer Science	Jasmin	Jones
7	French	Alice	Wilkins
6	English	Carol	Miller
3	French	Judy	Miller

(8 rows)

General Query

About
Cassandra
Differences
Key Features
Tables and
Columns
SQL vs CQL
Start
Cassandra
Keyspaces

Inserting
Data

Creating
Indexes

Shutting
Down

Consider this

Simple Query

```
select * from emp where empid = 1;
```

```
cqlsh:mydb> select * from emp where empid = 1;
```

empid	emp_dept	emp_first	emp_last
1	English	Fred	Smith

(1 rows)

Simple Query

```
create index idx_first on emp(emp_first); /* run this once, ever in DB */
select * from emp where emp_first = 'Fred';
```

```
cqlsh:mydb> create index idx_first on emp(emp_first);
cqlsh:mydb> select * from emp where emp_first = 'Fred';
```

empid	emp_dept	emp_first	emp_last
1	English	Fred	Smith

(1 rows)

Creating Indexes

- In Cassandra, if you want to query columns other than the primary key, you need to create a secondary index on them

```
select * from emp where emp_dept = 'English';
```

Try creating an index and run again...

```
create index idx_dept on emp(emp_dept);  
select * from emp where emp_dept = 'English';
```

```
select * from emp where emp_first = 'Fred';
```

```
create index idx_first on emp(emp_first);  
select * from emp where emp_first = 'Fred';
```

```
select * from emp where emp_last = 'Miller';
```

```
create index idx_last on emp(emp_last);  
select * from emp where emp_last = 'Miller';
```

More complex queries

Insert more data

```
insert into emp
(empid, emp_first, emp_last, emp_dept)
values (5, 'Megan', 'Douglas', 'French');
insert into emp
(empid, emp_first, emp_last, emp_dept)
values (6, 'Carol', 'Miller', 'English');
insert into emp
(empid, emp_first, emp_last, emp_dept)
values (7, 'Alice', 'Wilkins', 'French');
insert into emp
(empid, emp_first, emp_last, emp_dept)
values (8, 'Monroe', 'Monderson', 'Computer Science');
```

To do a more complicated query, we first have to index the column.

```
/* Create an index to find a last name */
create index idx_last on emp(emp_last);
describe index idx_last /*a working index?*/
describe index idx_first /*a working index?*/

select * from emp where emp_last = 'Alison';
```

How to Shut Down a Session

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Remove a keyspace; destroy data

```
/*Drop the "mydb" keyspace*/  
DROP KEYSPACE mydb;
```

Remove a table; destroy data

```
/*Drop the "emp" table*/  
DROP TABLE emp;
```

Closing down

- exit in the client terminal
- Control-C in the server terminal

Consider this...

About
Cassandra

Differences

Key Features

Tables and
Columns

SQL vs CQL

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



THINK

- Can you create and populate a new Cassandra database?