



Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Introduction to Database Systems: CS312

Cassandra

Oliver Bonham-Carter

10 April 2019

Quiz 2

Let's Review a Bit

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



Time to Review

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Theory of NoSQL

- The CAP Theorem
- The theory behind MongoDB and Cassandra
- Shards and scaling in MongoDB
- Schemas: NoSQL versus RDBMS

Compare and contrast

- Tools: Sqlite3 versus MongoDB
- Database Language Guide: terminologies of RDBMS compared to those of MongoDB
- *Sharding* and *scaling* in MongoDB
- Schemas: MongoDB versus Sqlite3
- The CAP theory (Capacity, Availability and Partition tolerance)

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Simple answers

- Provide short answers or basic code. (Review your sandbox files!)
 - Placing simple data into JSON format (i.e., a person's "name", "age", etc)
 - Inserting json data in MongoDB
 - Queries in MongoDB (see the *employee* data and its queries)
 - Comparing values in MongoDB queries
 - Queries in Cassandra

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Nodes

- All the nodes in a cluster play the same role.
 - Each node is independent and at the same time interconnected to other nodes.
- Each node in a cluster can accept read and write requests, regardless of where the data is actually located in the cluster.
- When a node goes down, read/write requests can be served from other nodes in the network.

Replication of Data

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Nodes

- One or more of the nodes in a cluster act as *replicas* (i.e., cloning) for a given piece of data
- If out-of-date values return from a node, Cassandra will return the most recent value to the client.
- After returning the most recent value, Cassandra performs a read repair in the background to update the stale values.

Replication of Data

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms


Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



Nodes

- Nodes are updated in the background using a *Gossip Protocol*
- Nodes communicate: detection of faulty nodes within cluster

Request Coordination

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms


Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



- Coordinator: the node chosen by the client to receive a particular read or write to its cluster
- No single point of Failure

Request Coordination

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms


Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



- The client connects to the closest node to run a query
- Each client request may be coordinated by a different node

Request Coordination

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms


Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



- The request is immediately sent to all other nodes where updates occur (if a write) or the query is performed (if a read)

Request Coordination

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms


Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this



- By the node responses, the coordinator node may also learn of any connectivity issues.

Request Coordination

Review

Data Replication

Request Coordination

Cassandra Query Language

Terms

Start Cassandra

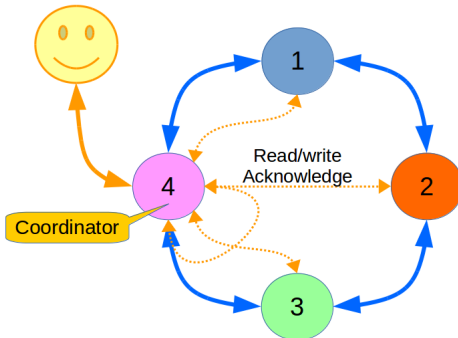
Keyspaces

Inserting Data

Shutting Down


Consider this

Client: **Cassandra Driver**



- The Cassandra driver chooses the coordinator node
- Round-Robin pattern, token-aware pattern: nodes attempt contact with others using different connection attempts each time and service is received from different nodes each time to help distribute load on the network.

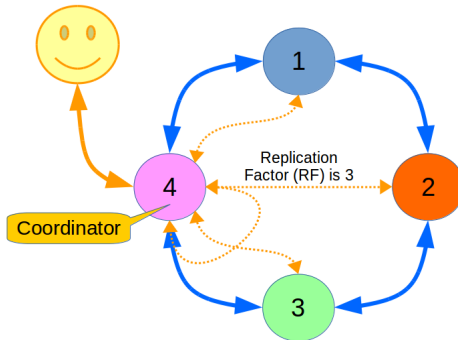
Round Robin?



- Signalling nodes use different orders to contact others each time to distribute the load across network

Request Coordination

Client: **Cassandra Driver**



- The coordinator manages the replication process
- Replication Factor (RF): num of nodes to share with
- The write will occur on the nodes responsible for that partition
- Each write has been time-stamped

Access to Cassandra

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

cqlsh

- Access Cassandra through its nodes using Cassandra Query Language (CQL).
- CQL treats the database (Keyspace) as a container of tables.
- Note: Programmers use `cqlsh`: a prompt to work with CQL or use separate application language drivers to access the base.
- Clients approach any of the nodes for their read-write operations.
- That node (coordinator) plays a proxy between the client and the nodes holding the data.

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this


Read Operations

During read operations, Cassandra gets values from the mem-table and checks the bloom filter to find the appropriate SSTable that holds the required data.

Write Operations

- Write activities of nodes are captured by the commit logs written in the nodes.
- This data will eventually be captured and stored in the mem-table.
- Whenever the mem-table is full, data will be written into the SSTable data file.
- All writes are automatically partitioned and replicated throughout the cluster.
- Garbage Collection: Cassandra periodically consolidates the SSTables, discarding unnecessary data.

Terms and Definitions



Term	Definition
Node	The place where data is stored.
Data center	A collection of related nodes.
Cluster	A cluster is a component that contains one or more data centers.
Commit log	A crash-recovery mechanism in Cassandra. Every write operation is written to the commit log to recover in case of a crash.

Terms and Definitions

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Term	Definition
Mem-table	A mem-table is a memory-resident data structure. After commit is logged, the data will be written to the mem-table.
SSTable	A disk file to which the data is flushed from the mem-table when its contents reach a threshold value.
Bloom filter	A collection of quick, nondeterministic, algorithms for testing whether an element is a member of a set. A cache; bloom filters are accessed after every query.

Getting started with Cassandra

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

- Copy your Cassandra setup tar file (apache-cassandra-3.11.2-bin.tar.gz) to a desktop directory. Do not copy this file or open it in your submission directory!!!
- Click on this file to unpack its contents
- Locate the bin directory and then locate file: `cassandra` and file: `cqlsh`

Lab PCs: Setup Java 8 in the path from the terminal

```
JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/
```

Start the Cassandra server

Note: Control-C to exit.

```
./cassandra -f
```

With new another terminal, start the `cqlsh` client

```
./cqlsh
```

Keyspaces

Similar to a schema

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Checking your database

```
describe keyspaces;  
describe schema;  
describe tables;  
describe table tableName;
```

Start a new keyspaces

```
create keyspace mydb with replication =  
{ 'class':'SimpleStrategy', 'replication_factor':1 };
```

Use a keyspace

```
use mydb;
```

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Build Tables

```
CREATE TABLE songs (  
    id uuid PRIMARY KEY,  
    title text,  
    album text,  
    artist text,  
    data blob  
);
```

```
CREATE TABLE playlists (  
    id uuid,  
    song_order int,  
    song_id uuid,  
    title text,  
    album text,  
    artist text,  
    PRIMARY KEY (id, song_order ) );
```

```
describe tables;
```

Example Music DB

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Enter playlists Data

```
INSERT INTO playlists (id, song_order, song_id, title, artist, album)
VALUES (62c36092-82a1-3a00-93d1-46196ee77204, 4,
7db1a490-5878-11e2-bcfd-0800200c9a66,
'Ojo Rojo', 'Fu Manchu', 'No One Rides for Free');
```

```
INSERT INTO playlists (id, song_order, song_id, title, artist, album)
VALUES (62c36092-82a1-3a00-93d1-46196ee77206, 3,
7db1a490-5878-11e2-bcfd-0800200c9a68,
'Green and The Other Colors in the Rainbow', 'Kermit the Frog', 'Funky Lilly Pad');
```

```
INSERT INTO playlists (id, song_order, song_id, title, artist, album)
VALUES (62c36092-82a1-3a00-93d1-46196ee77207, 1,
7db1a490-5878-11e2-bcfd-0800200c9a69,
'Shovin the Lovin -- cookies in my mouth', 'Cookie Monster', 'Blue Carpet Disaster');
```

```
INSERT INTO playlists (id, song_order, song_id, title, artist, album)
VALUES (62c36092-82a1-3a00-93d1-46196ee77205, 2,
7db1a490-5878-11e2-bcfd-0800200c9a67,
'Singin With Lovin', 'Roy Ray Robinson', 'The Sun Shine on Your Shoulders');
```

```
INSERT INTO playlists (id, song_order, song_id, title, artist, album)
VALUES (62c36092-82a1-3a00-93d1-46196ee77207, 1,
7db1a490-5878-11e2-bcfd-0800200c9a69,
'Dance to me', 'Dance Floor Dolls', 'Electric Motion');
```

Example Music DB

We will only enter one row, *just for show*

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Enter songs Data

```
INSERT INTO songs (id, title, album, artist, data)
VALUES
  (62c36092-82a1-3a00-93d1-46196ee77204,
   'Fu_title',
   'Fu_album',
   'Fu_artist',
   textAsBlob('Fu_blob'));
```

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Simple queries

```
select * from playlists;  
select album from playlists;
```

Simple queries with conditions

```
select * from playlists  
  where id = 62c36092-82a1-3a00-93d1-46196ee77205;  
  
select album from playlists  
  where id = 62c36092-82a1-3a00-93d1-46196ee77205;  
  
select  album, title from playlists  
  where id = 62c36092-82a1-3a00-93d1-46196ee77205;
```


Conditional Queries

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

- In Cassandra, if you want to query columns other than the primary key, you need to create a secondary index on them

Simple queries with conditions

```
CREATE INDEX ON playlists( artist );  
/*Now, you can query the playlists for songs by Fu Manchu. */
```

Index membership code

```
select * from playlists where artist = 'Fu Manchu';  
select * from playlists where artist = 'Kermit the Frog';
```

Study the schema and index memberships (i.e., *query-able* columns)

```
describe schema
```

How to Shut Down a Session

Review

Data
Replication

Request
Coordination

Cassandra
Query
Language

Terms

Start
Cassandra

Keyspaces

Inserting
Data

Shutting
Down

Consider this

Remove a keyspace; destroy data

```
/*Drop the "mydb" keyspace*/  
DROP KEYSPACE mydb;
```

Remove a table; destroy data

```
/*Drop the "emp" table*/  
DROP TABLE emp;
```

Closing down

- exit in the client terminal
- Control-C in the server terminal

Consider this...



THINK

- Can you query other columns in the `playlists` table?
- Can you write index membership code to allow for the query of other columns?